

ABSTRACT

This project focuses on optimizing computationally intensive functions using ARM NEON intrinsics for embedded auto-control systems. NEON is ARM's SIMD (Single Instruction, Multiple Data) extension that enables parallel processing. The primary goal was to accelerate functions such as histogram calculation, grayscale conversion, Sobel edge detection, and box blurring. These are crucial in real-time perception systems like ADAS (Advanced Driver Assistance Systems).

NEON-based optimizations significantly reduced processing time across all functions. For instance, grayscale conversion showed 2 \times speedup, and Sobel filtering showed a 3 \times improvement. These results demonstrate that NEON intrinsics can greatly enhance system responsiveness, lower power consumption, and improve performance in embedded vision applications.

Contents

1	Introduction	8
1.1	Motivation	9
1.2	Objectives	9
1.3	Problem statement	10
1.4	Application in Real-Life Context.....	10
1.5	Project Planning and Bill of Materials.....	10
1.5.1	Project Planning	10
1.5.2	Bill of Material.....	11
1.6	Organization of the report	11
2	System Design	12
2.1	Functional block Diagram	12
2.2	Design Alternatives.....	13
2.3	Final Design	13
3	Implementation Details	14
3.1	Specifications and System Architecture.....	14
3.2	Software Stack Details.....	15
3.2.1	Histogram calculations	15
3.2.2	Grayscale Conversions.....	18
3.2.3	Sobel filter for edge detection	19
3.2.4	Box blurring	20
3.3	Algorithm.....	21
3.4	Flowchart.....	22
4	Results and Discussions	23
4.1	Performance analysis.....	23
5	Conclusions and Future Scope	25
5.1	Conclusion	25
5.2	Future Scope	25

List of Tables

1.1 Bill of Materials	11
4.1 Performance analysis	23

List of Figures

2.1	Block diagram depicting overview of the system operation.....	12
3.1	Linear data loading.....	16
3.2	Structural data loading.....	16
3.3	Before optimization	17
3.4	After optimization	17
3.5	RGB image	18
3.6	Gray Scale image	19
3.7	Blurred image.....	19
3.8	Sobel filter	20
3.9	Edge detection.....	20
3.10	Flowchart.....	22
4.1	Statistical block data	23
4.2	Histogram.....	24
4.3	Snapshot of terminal.....	24

Chapter 1

Introduction

In modern embedded systems, particularly within automotive electronics, the ability to process visual data in real time is critical for applications such as Adaptive Cruise Control, Lane Keeping Assist, Obstacle Detection, and Driver Monitoring Systems. These applications demand high-performance image processing under strict latency and power constraints, especially on edge-AI platforms like the NVIDIA Jetson AGX Orin.

Traditional scalar implementations of image processing algorithms—such as grayscale conversion, histogram analysis, edge detection, and blurring—often become performance bottlenecks due to the sheer volume of pixel data and the need for real-time responsiveness. To address these limitations, this project proposes a function-level optimization framework leveraging ARM's NEON SIMD (Single Instruction, Multiple Data) architecture.

The project focuses on rewriting critical image processing functions using NEON intrinsics, enabling data-level parallelism to dramatically accelerate execution on Jetson's ARM Cortex-A78AE CPUs. NEON provides 128-bit vector registers that allow simultaneous processing of multiple pixels or operations in a single instruction cycle, significantly improving throughput without increasing CPU frequency or power consumption.

The optimized system is benchmarked against baseline (non-NEON) implementations for four key functions: grayscale conversion, Sobel filter-based edge detection, histogram calculation, and box blurring. Each function was carefully profiled and restructured using NEON load-store operations, vectorized arithmetic, and approximation techniques (e.g., bit-trick for exponential functions) to demonstrate real-world performance gains.

The results highlight substantial improvements, such as a 2x speedup in grayscale conversion and a 3x reduction in edge detection latency. These optimizations directly translate into better frame rates and reduced energy consumption—essential for real-time ADAS systems deployed in resource-constrained environments.

The modular design of this optimization framework and its compatibility with common edge-AI platforms make it highly reusable for a wide range of vision-based applications. Future extensions could include deep learning model acceleration, GPU-CPU hybrid optimization, or integration with real-time sensor fusion pipelines for complete ADAS solutions.

1.1 Motivation

Embedded systems for ADAS and autonomous driving require real-time image processing with minimal latency. Conventional scalar code struggles under these constraints. NEON intrinsics allow multiple data points to be processed simultaneously, making them ideal for speeding up image preprocessing tasks like histogram calculation and edge detection.

1.2 Objectives

- To identify and profile key image processing functions in automotive control systems that act as performance bottlenecks (e.g., grayscale conversion, histogram calculation, edge detection, and box blurring).
- To implement NEON-optimized versions of these functions using ARM SIMD intrinsics on the NVIDIA Jetson AGX Orin platform.
- To leverage 128-bit NEON vector registers and parallel data processing techniques to achieve significant reductions in execution time and latency.
- To conduct side-by-side benchmarking and performance analysis between baseline (scalar) and NEON-optimized implementations across 1000+ iterations.
- To incorporate additional mathematical optimizations, such as bit-trick approximations for exponential operations used in histogram compression.
- To demonstrate the functional correctness and real-time gains of NEON-optimized code through quantifiable metrics like speedup factor, execution time (μs), and resource utilization.
- To develop a modular, reusable, and low-cost software optimization framework for accelerating vision-based embedded systems in edge AI, ADAS, and robotics.

1.3 Problem statement

Neon Optimization of Function from Auto Control

Real-time functions in ADAS like lane detection, object tracking, and sensor fusion require fast, low-power processing. Without optimization, these algorithms risk missing deadlines or overloading embedded CPUs. This project aims to close that gap using NEON intrinsics.

1.4 Application in Real-Life Context

- Adaptive Cruise Control and Lane Keeping
- Real-time video processing in drones and robotics
- Surveillance systems with edge AI processing
- Smart traffic monitoring and autonomous navigation

1.5 Project Planning and Bill of Materials

1.5.1 Project Planning

This project focuses on optimizing core image-processing functions using ARM NEON intrinsics on the Jetson AGX Orin platform. These functions—such as histogram calculation, grayscale conversion, and edge detection—are essential in real-time auto-control systems. The workflow includes NEON-based optimization, benchmarking, and performance comparison with baseline implementations to validate speed and efficiency improvements.

System Setup:

- Jetson AGX Orin with a computer.

Data Validation:

- Performance of each optimized function is benchmarked against its baseline implementation.
- Speedup and correctness are verified using execution time analysis across multiple iterations.

1.5.2 Bill of Material

Component	Description	Quantity	Cost (INR)
Jetson AGX Orin Dev Kit	Controller Device (CPU)	1	2,50,000

Table 1.1: Bill of Material

1.6 Organization of the report

This report is organized into five chapters, each covering a distinct aspect of the I2C Camera Emulator project. Below is an overview of the chapter structure:

- **Chapter 2 – System Design:** This chapter describes the overall system architecture, including the target hardware platform (Jetson AGX Orin), the role of NEON SIMD optimization, and the selection of functions chosen for acceleration.
- **Chapter 3 – Implementation Details:** This chapter provides the implementation specifics for each optimized function—such as histogram calculation, grayscale conversion, Sobel filtering, and box blurring. It includes the use of ARM NEON intrinsics, code-level techniques, and algorithmic strategies.
- **Chapter 4 – Results and Discussions:** This chapter presents performance benchmarks comparing NEON-optimized functions with their baseline versions. It includes speedup factors, execution time analysis, and observations on computational efficiency.
- **Chapter 5 – Conclusions and Future Scope:** The final chapter summarizes key outcomes, emphasizes the impact of NEON optimization, and discusses potential extensions such as support for higher resolution data, GPU-based optimization, or real-time deployment in ADAS systems.

Chapter 2

System Design

The **Neon Optimization framework** is designed to accelerate critical image-processing functions in embedded auto-control systems using ARM's NEON SIMD architecture. Implemented on the NVIDIA Jetson AGX Orin platform, this project targets real-time performance bottlenecks in operations like histogram calculation, grayscale conversion, edge detection, and box blurring. By leveraging NEON intrinsics, the system enables high-throughput, low-latency processing suitable for applications such as ADAS. This setup provides a reproducible and performance-focused environment for evaluating the computational efficiency of optimized functions, allowing developers to enhance the responsiveness and reliability of embedded vision pipelines without the need for additional hardware or external accelerators.

2.1 Functional Block Diagram




Figure 2.1: Block diagram depicting overview of the system operation.

2.2 Design Alternatives

In the development of the **Neon Optimization framework**, several design alternatives were explored for optimizing key image-processing functions before finalizing the current implementation. These alternatives focused on **optimization strategies, instruction-level techniques, and performance validation methods**, all evaluated based on execution speed, resource efficiency, and compatibility with the Jetson AGX Orin platform.

Key alternatives considered:

- **Manual Loop Unrolling vs NEON Intrinsics:**
Loop unrolling was initially considered to reduce loop overhead, but NEON intrinsics provided better parallelism and scalability with SIMD registers.
- **Taylor Series vs Bit-Trick for Exponential Calculations:**
For functions involving exponential compression (e.g., histogram normalization), bit-level tricks were preferred over Taylor expansion due to superior execution speed and minimal memory overhead.
- **Inline Assembly vs Compiler Intrinsics:**
While inline assembly offered finer control, NEON compiler intrinsics were chosen for better readability, portability, and ease of integration into existing C/C++ codebases.
- **Function-specific vs Generic SIMD Modules:**
Initially, a generic SIMD template was considered for reuse across functions, but dedicated NEON routines were found to be more efficient due to their tailored memory access patterns.

2.3 Final Design

After evaluating various approaches, the following design decisions were adopted to maximize performance and maintain code clarity:

- **NEON Intrinsics:**
ARM NEON intrinsics were selected over inline assembly for optimizing image-processing functions due to their portability, ease of use, and effective parallelism.
- **Bit-Trick Optimization:**
For exponential operations (used in histogram normalization), bit-trick methods were chosen for faster execution over traditional Taylor series expansion.
- **Dedicated SIMD Paths:**
Instead of using a generic SIMD module, each function was optimized individually to take advantage of specific data patterns, improving efficiency.
- **Function-Specific Benchmarks:**
Each function was profiled and benchmarked independently to accurately measure performance improvements and avoid cross-function interference.

These choices enabled a clean, modular implementation that delivered significant speedups while remaining scalable for embedded deployment on Jetson AGX Orin.

Chapter 3

Implementation Details

3.1 Specifications and System Architecture

The Neon Optimization framework is developed and tested on the NVIDIA Jetson AGX Orin platform to accelerate performance-critical functions in embedded vision systems. These functions include histogram calculation, grayscale conversion, Sobel edge detection, and box blurring. Each function is optimized using ARM NEON intrinsics to enable parallel processing and reduce execution time.

This system follows a modular software architecture where each function is profiled, optimized, and benchmarked independently. The design ensures flexibility, allowing easy integration of future functions and scalable performance analysis.

Hardware Specifications:

- **Jetson AGX Orin Dev Kit:**
High-performance ARM Cortex-A78AE-based SoC with 12 CPU cores at 2.2 GHz, NEON SIMD support, and up to 64GB LPDDR5 memory.

Software Stack:

- C/C++ source code:
Implements both baseline and NEON-optimized versions of each function using ARM intrinsics.
- GCC Compiler with NEON Flags:
Uses -mfpu=neon or -march=armv8-a+simd flags for building optimized binaries.
- Benchmarking Scripts:
Measure execution time and speedup for 1000 iterations of each function.
- Terminal/Console Output:
Logs performance metrics, function outputs, and validation results.

3.2 Software Stack Details

The software stack for the Neon Optimization project consists of tightly integrated components developed in C/C++ and executed on the Jetson AGX Orin platform. The primary goal is to implement, optimize, and benchmark various image-processing functions using ARM NEON intrinsics.

Each function—such as histogram computation, grayscale conversion, Sobel edge detection, and box blurring—is written in two variants: a baseline scalar version and a NEON-optimized version using SIMD instructions. These are compiled with NEON support enabled and executed on the Jetson CPU.

For performance analysis, each function is profiled across 1000 iterations, and metrics such as execution time and speedup factor are logged. The results help compare the efficiency of NEON optimizations against traditional implementations.


Additionally, terminal outputs and optional log files are used to validate functional correctness and summarize benchmarking data for analysis

3.2.1 Histogram calculations

Histogram Calculations “`CalcColorHistogramsFromStatsBlock()`”


- Loading RGB averages and weights in parallel.
- Computing the maximum color intensity per window.
- Compressing values (like squeezing a range) using custom power logic.
- Rounding and clamping indices for histogram bins.
- Updating the histogram by accumulating weights based on intensity.

The NEON SIMD pipeline replaces 4 iterations of scalar math with 1 instruction, giving major speedups.



Loading RGB data with a linear load.

Figure 3.1: Linear data loading



Loading RGB data with a structure load.

Figure 3.2: Structural data loading

```

for (UInt32 i = 0; i < ROI_NUM; i++)
{
    if (ROIEnable[i])
    {
        for (UInt32 j = 0; j < NumWindows[i]; j++)
        {
            // Inlined implementation of GetWeightTableIndexForStatsBlock
            UInt32 idx;
            if (roiWidth == 0)
            {
                idx = 0;
            }
            else
            {
                UInt32 wtIndexForRoi = roiWidth * (i % 2) + roiWidth * wtWidth * (i / 2);
                idx = wtIndexForRoi + j % roiWidth + wtWidth * (j / roiWidth);
            }

            const Float32 weight = weightTablePtr[idx];
            const Float32 r = pAverage_R[i * MAX_ROI_WINDOWS + j];
            const Float32 g = pAverage_G[i * MAX_ROI_WINDOWS + j];
            const Float32 b = pAverage_B[i * MAX_ROI_WINDOWS + j];
            const Float32 maxValueFloat =
                static_cast<Float32>(powf(MAX_OF_THREE(r,g,b), histogramCompressionPower) *
                (HISTOGRAM_SIZE - 1));

            const UInt32 maxValueInt = static_cast<UInt32>(maxValueFloat + 0.5f);
            if (maxValueInt > (HISTOGRAM_SIZE - 1))
            {
                printf("Calculated histogram index exceeds limit\n");
            }

            histPtrOutput[maxValueInt] += weight;
            totalMass += weight;
        }
    }
}

```

Figure 3.3: Before optimization

```

// Process ROIs and accumulate histogram values
for (UInt32 i = 0; i < ROI_NUM; i++) {
    if (ROIEnable[i]) {
        for (UInt32 j = 0; j < NumWindows[i]; j += 4) { // Process in chunks of 4 (vectorized)
            UInt32 idx[4];
            for (int k = 0; k < 4; k++) {
                UInt32 wtIndexForRoi = roiWidth * (i % 2) + roiWidth * wtWidth * (i / 2);
                idx[k] = wtIndexForRoi + (j + k) % roiWidth + wtWidth * ((j + k) / roiWidth);
            }

            float32x4_t weights = vld1q_f32(&weightTablePtr[idx[0]]);
            float32x4_t r = vld1q_f32(&pAverage_R[i * MAX_ROI_WINDOWS + j]);
            float32x4_t g = vld1q_f32(&pAverage_G[i * MAX_ROI_WINDOWS + j]);
            float32x4_t b = vld1q_f32(&pAverage_B[i * MAX_ROI_WINDOWS + j]);

            // Compute maximum of r, g, b
            float32x4_t max_rg = vmaxq_f32(r, g);
            float32x4_t max_rgb = vmaxq_f32(max_rg, b);

            // Apply compression power
            float32x4_t max_value_float = vmulq_n_f32(
                vector_pow(max_rgb, histogramCompressionPower),
                (HISTOGRAM_SIZE - 1));
        }

        // Round to nearest integer and clamp
        uint32x4_t max_value_int = vcvtq_u32_f32(
            vaddq_f32(max_value_float, vdupq_n_f32(0.5f)));
        uint32x4_t clamped_indices = vminq_u32(
            max_value_int,
            vdupq_n_u32(HISTOGRAM_SIZE - 1));
    }

    // Update histogram bins
    for (int k = 0; k < 4; k++) {
        UInt32 index = vgetq_lane_u32(clamped_indices, k);
        histPtrOutput[index] += vgetq_lane_f32(weights, k);
        totalMass += vgetq_lane_f32(weights, k);
    }
}

```

Figure 3.4: After optimization

- All the possible operations involving load and store were optimized using **NEON** Intrinsics
- Loading and storing intrinsics:

```
vld1q_f32(&array[index])           //Loads 4 floats at once

vst1q_f32(result, [1.0, 2.0, 3.0, 4.0]); //Stores 4 floats at once
```

3.2.2 GrayScale Conversion

Convert RGB pixels to grayscale using:

$$\text{GRAY} = 0.299 \cdot R + 0.587 \cdot G + 0.114 \cdot B$$

Neon optimization algorithm:

- Preload Weights as NEON Vectors
- Process Pixels in Chunks of 8
- Multiply-Accumulate and Downscale
- Stores 8 grayscale pixels at once



Figure 3.5: RGB image



Figure 3.6: Gray scale image

3.2.3 Box blurring

A box blur is a simple image filter where each output pixel is the average of a surrounding box of pixels in the input image.

The optimized function applies a box blur using ARM NEON SIMD intrinsics on grayscale image data (uint8_t values), processing 16 pixels at once horizontally.




Figure 3.7: Blurred image

3.2.4 Sobel filter for edge detection

- When we apply this mask on the image it prominent vertical edges. It simply works like as first order derivative and calculates the difference of pixel intensities in a edge region.




Figure 3.8: Sobel Filter

SobelNeon Function:

- It processes 8 pixels in parallel using NEON Registers (`uint8x8_t`)
- Each 3x3 neighbourhood around a pixel is loaded using `vld1_u8`



Figure 3.9: Edge detection

3.3 Algorithm

Jetson AGX:

- **Select the target image-processing function**
(e.g., Histogram, Grayscale Conversion, Sobel Filter, or Box Blur).
- **Load the input data**
(e.g., pixel array or statistical block) into memory.
- **Execute the baseline (non-NEON) version**
and record the execution time.
- **Execute the NEON-optimized version**
using ARM intrinsics with SIMD instructions.
- **Benchmark and compare**
both implementations by calculating speedup:
 $\text{Speedup} = \text{Baseline Time} / \text{NEON Time}$.
- **Validate output accuracy**
by comparing outputs from both versions for correctness.
- **Log the performance data**
(execution time, speedup, memory usage, etc.) for reporting.

3.4 Flowchart




Figure 3.10: Flowchart

Chapter 4

Results and Discussions

4.1 Performance analysis

To evaluate the performance impact of NEON optimization, each target function—Histogram Calculation, Grayscale Conversion, Box Blurring, and Sobel Edge Detection—was implemented in both scalar (baseline) and NEON-intrinsic versions. The benchmarking was performed on the NVIDIA Jetson AGX Orin, which features a 12-core ARM Cortex-A78AE CPU with NEON SIMD support.

Each function was executed over **1000 iterations**, and the execution times were averaged for both the baseline and NEON-optimized versions. The **speedup factor** was calculated by comparing these two execution times, giving a quantitative measure of performance improvement.

Function	Baseline time (μs)	NEON Time (μs)	Speedup Factor
Histogram	127	86	1.5x
Grayscale	2919	1451	~2x
Box blurring	1040.73	287.49	~3.6x
Edge detection	19723	6592	~3x

Table 4.1: Performance analysis

```
testcase 1: No Compression
#0.003906 Bin 1: 0.003906 Bin 2: 0.003906 Bin 3: 0.003906 Bin 4: 0.003906 Bin 5: 0.003906
Bin 6: 0.003906 Bin 7: 0.003906 Bin 8: 0.003906 Bin 9: 0.003906 Bin 10: 0.003906 Bin 11: 0.003906 Bin 12: 0.003906
Bin 13: 0.003906 Bin 14: 0.003906 Bin 15: 0.003906 Bin 16: 0.003906 Bin 17: 0.003906 Bin 18: 0.003906 Bin 19: 0.003906
Bin 20: 0.003906 Bin 21: 0.003906 Bin 22: 0.003906 Bin 23: 0.003906 Bin 24: 0.003906 Bin 25: 0.003906
Bin 26: 0.003906 Bin 27: 0.003906 Bin 28: 0.003906 Bin 29: 0.003906 Bin 30: 0.003906 Bin 31: 0.003906 Bin 32: 0.003906 Bin 33: 0.003906
Bin 34: 0.003906 Bin 35: 0.003906 Bin 36: 0.003906 Bin 37: 0.003906 Bin 38: 0.003906 Bin 39: 0.003906 Bin 40: 0.003906
Bin 41: 0.003906 Bin 42: 0.003906 Bin 43: 0.003906 Bin 44: 0.003906 Bin 45: 0.003906 Bin 46: 0.003906
Bin 47: 0.003906 Bin 48: 0.003906 Bin 49: 0.003906 Bin 50: 0.003906 Bin 51: 0.003906 Bin 52: 0.003906 Bin 53: 0.003906 Bin 54: 0.003906
Bin 55: 0.003906 Bin 56: 0.003906 Bin 57: 0.003906 Bin 58: 0.003906 Bin 59: 0.003906 Bin 60: 0.003906
Bin 61: 0.003906 Bin 62: 0.003906 Bin 63: 0.003906 Bin 64: 0.003906 Bin 65: 0.003906 Bin 66: 0.003906
Bin 67: 0.003906 Bin 68: 0.003906 Bin 69: 0.003906 Bin 70: 0.003906 Bin 71: 0.003906 Bin 72: 0.003906
Bin 73: 0.003906 Bin 74: 0.003906 Bin 75: 0.003906 Bin 76: 0.003906 Bin 77: 0.003906
Bin 78: 0.003906 Bin 79: 0.003906 Bin 80: 0.003906 Bin 81: 0.003906 Bin 82: 0.003906 Bin 83: 0.003906
Bin 84: 0.003906 Bin 85: 0.003906 Bin 86: 0.003906 Bin 87: 0.003906 Bin 88: 0.003906 Bin 89: 0.003906
Bin 90: 0.003906 Bin 91: 0.003906 Bin 92: 0.003906 Bin 93: 0.003906 Bin 94: 0.003906 Bin 95: 0.003906
Bin 96: 0.003906 Bin 97: 0.003906 Bin 98: 0.003906 Bin 99: 0.003906 Bin 100: 0.003906 Bin 101: 0.003906
Bin 102: 0.003906 Bin 103: 0.003906 Bin 104: 0.003906 Bin 105: 0.003906 Bin 106: 0.003906 Bin 107: 0.003906
Bin 108: 0.003906 Bin 109: 0.003906 Bin 110: 0.003906 Bin 111: 0.003906 Bin 112: 0.003906 Bin 113: 0.003906
Bin 114: 0.003906 Bin 115: 0.003906 Bin 116: 0.003906 Bin 117: 0.003906 Bin 118: 0.003906 Bin 119: 0.003906
Bin 120: 0.003906 Bin 121: 0.003906 Bin 122: 0.003906 Bin 123: 0.003906 Bin 124: 0.003906 Bin 125: 0.003906
Bin 126: 0.003906 Bin 127: 0.003906 Bin 128: 0.003906 Bin 129: 0.003906 Bin 130: 0.003906 Bin 131: 0.003906
Bin 132: 0.003906 Bin 133: 0.003906 Bin 134: 0.003906 Bin 135: 0.003906 Bin 136: 0.003906 Bin 137: 0.003906
Bin 138: 0.003906 Bin 139: 0.003906 Bin 140: 0.003906 Bin 141: 0.003906 Bin 142: 0.003906 Bin 143: 0.003906
Bin 144: 0.003906 Bin 145: 0.003906 Bin 146: 0.003906 Bin 147: 0.003906 Bin 148: 0.003906 Bin 149: 0.003906
Bin 150: 0.003906 Bin 151: 0.003906 Bin 152: 0.003906 Bin 153: 0.003906 Bin 154: 0.003906 Bin 155: 0.003906
Bin 156: 0.003906 Bin 157: 0.003906 Bin 158: 0.003906 Bin 159: 0.003906 Bin 160: 0.003906 Bin 161: 0.003906
Bin 162: 0.003906 Bin 163: 0.003906 Bin 164: 0.003906 Bin 165: 0.003906 Bin 166: 0.003906 Bin 167: 0.003906
Bin 168: 0.003906 Bin 169: 0.003906 Bin 170: 0.003906 Bin 171: 0.003906 Bin 172: 0.003906 Bin 173: 0.003906
Bin 174: 0.003906 Bin 175: 0.003906 Bin 176: 0.003906 Bin 177: 0.003906 Bin 178: 0.003906 Bin 179: 0.003906
Bin 180: 0.003906 Bin 181: 0.003906 Bin 182: 0.003906 Bin 183: 0.003906 Bin 184: 0.003906 Bin 185: 0.003906
Bin 186: 0.003906 Bin 187: 0.003906 Bin 188: 0.003906 Bin 189: 0.003906 Bin 190: 0.003906 Bin 191: 0.003906
Bin 192: 0.003906 Bin 193: 0.003906 Bin 194: 0.003906 Bin 195: 0.003906 Bin 196: 0.003906 Bin 197: 0.003906
Bin 198: 0.003906 Bin 199: 0.003906 Bin 200: 0.003906 Bin 201: 0.003906 Bin 202: 0.003906 Bin 203: 0.003906
Bin 204: 0.003906 Bin 205: 0.003906 Bin 206: 0.003906 Bin 207: 0.003906 Bin 208: 0.003906 Bin 209: 0.003906
Bin 210: 0.003906 Bin 211: 0.003906 Bin 212: 0.003906 Bin 213: 0.003906 Bin 214: 0.003906 Bin 215: 0.003906
Bin 216: 0.003906 Bin 217: 0.003906 Bin 218: 0.003906 Bin 219: 0.003906 Bin 220: 0.003906 Bin 221: 0.003906
Bin 222: 0.003906 Bin 223: 0.003906 Bin 224: 0.003906 Bin 225: 0.003906 Bin 226: 0.003906 Bin 227: 0.003906
Bin 228: 0.003906 Bin 229: 0.003906 Bin 230: 0.003906 Bin 231: 0.003906 Bin 232: 0.003906 Bin 233: 0.003906
Bin 234: 0.003906 Bin 235: 0.003906 Bin 236: 0.003906 Bin 237: 0.003906 Bin 238: 0.003906 Bin 239: 0.003906
Bin 240: 0.003906 Bin 241: 0.003906 Bin 242: 0.003906 Bin 243: 0.003906 Bin 244: 0.003906 Bin 245: 0.003906
Bin 246: 0.003906
```

Figure 4.1: Statistical block data




Figure 4.2: Histogram

```

nvidia@ubuntu:~/Desktop/Neon optimization$ g++ -O3 test_code.cpp autocontrol_utils.cpp -lm -o normalcode
nvidia@ubuntu:~/Desktop/Neon optimization$ ./normalcode
average execution time: 127 microseconds
nvidia@ubuntu:~/Desktop/Neon optimization$ 

nvidia@ubuntu:~/Desktop/Neon optimization$ g++ -O3 -march=armv8-a+simd -std=c++17 -pg -g neon_opt.cpp -o neontime
nvidia@ubuntu:~/Desktop/Neon optimization$ ./neontimetest
Average execution time: 86 microseconds
nvidia@ubuntu:~/Desktop/Neon optimization$ 

```

Figure 4.3: A snapshot of terminal

Chapter 5

Conclusions and Future Scope

5.1 Conclusion

The Neon Optimization of Functions from Autocontrol project demonstrates the effectiveness of using ARM NEON intrinsics to accelerate core image-processing operations in embedded vision systems. Implemented on the NVIDIA Jetson AGX Orin platform, the project focused on optimizing functions such as histogram calculation, grayscale conversion, Sobel edge detection, and box blurring—each critical for real-time perception in applications like driver-assistance systems. By leveraging SIMD parallelism, significant reductions in execution time were achieved, validating the potential of NEON for enhancing performance under strict latency and power constraints. The optimized implementations maintain functional accuracy while offering substantial speedup over traditional scalar code. This project provides a scalable and practical framework for embedding high-efficiency vision algorithms in resource-constrained environments, establishing a strong foundation for future work in real-time embedded optimization.

5.2 Future Scope

While the current implementation fulfils its core objectives, several improvements and extensions are possible:

- **Dynamic Optimization Selection:** Integrate a configuration interface to toggle between baseline and NEON-optimized functions at runtime for adaptive performance evaluation.
- **Support for Additional Functions:** Expand the optimization framework to include other vision-based operations like median filtering, dilation, erosion, and thresholding.
- **Automated Benchmarking Suite:** Develop a GUI or script-based tool for auto-profiling, logging, and visualizing speedup and performance metrics across multiple functions.
- **Cross-Platform NEON Portability:** Extend the implementation to support other NEON-capable platforms (e.g., Raspberry Pi 4, i.MX series) for broader applicability.
- **Integration with Real-Time Systems:** Incorporate NEON-optimized modules into a live ADAS testbench or ROS-based robotic platform to evaluate real-world latency benefits.